


Lecture 8

Descriptive Statistics



Announcements

- ❖ Lab 2 today
 - Analyze statistics from a penguin dataset
 - Working in pairs
 - Improved lab process
- ❖ Friday starts Blue & Gold weekend
 - Rather than a normal lecture, you'll get a chance to meet with alumni and potential employers (probably)
 - You'll fill out an attendance card at the event

Last Few Times

- ❖ Using Python and Pandas
 - Loading data from csv
 - Filtering rows and columns
- ❖ GitHub for version control, storage, and distribution of projects

Today

- ❖ What is a statistic?
- ❖ Central tendency
- ❖ Spread and variation

What is a statistic?

- ❖ A summary of the data
- ❖ Often captures one aspect or feature of a dataset


What is a statistic?

- ❖ A summary of the data
- ❖ Often captures one aspect or feature of a dataset
- ❖ We often want to describe a dataset with a small number of statistics
- ❖ Numerical summaries
 - Averages, counts, quantiles, measures of spread, correlation coefficients
- ❖ Visualizations
 - Histograms, scatterplots, pie charts

Counts & Histograms

How many of each value are in our dataset?

- ❖ Gives us a full summary of the variable in our dataset




Number of Colors in Bag of M&M Candies

Counts & Histograms

How many of each value are in our dataset?

- ❖ Gives us a full summary of the variable in our dataset
- ❖ How might this go wrong?
 - Think about a numeric variable




Number of Colors in Bag of M&M Candies

Counts & Histograms

How many of each value are in our dataset?

- ❖ Gives us a full summary of the variable in our dataset
- ❖ How might this go wrong?
 - Think about a numeric variable
 - It's possible that every value is unique!




Number of Colors in Bag of M&M Candies

Counts & Histograms

How many of each value are in our dataset?

- ❖ Gives us a full summary of the variable in our dataset
- ❖ How might this go wrong?
 - Think about a numeric variable
 - It's possible that every value is unique!
- ❖ For numeric variables, we tend to "bin" the data into "buckets"
 - Think of this as creating a new variable that's easier to visualize from the existing one



Centrality

- ❖ For variables that can be ordered, it's helpful to know a typical value

$$x = [x_1, x_2, x_3, \dots, x_n]$$

Centrality

- ❖ For variables that can be ordered, it's helpful to know a typical value

$$X = [X_1, X_2, X_3, \dots, X_n]$$

- ❖ Mean (or average)
 - Sum all the values and divide by n ; denoted by $\bar{x}$
- ❖ Median
 - Sort the values and find the middle value (or average the two middle values)


Centrality

- ❖ For variables that can be ordered, it's helpful to know a typical value

$$X = [X_1, X_2, X_3, \dots, X_n]$$

- ❖ Mean (or average)
 - Sum all the values and divide by n ; denoted by $\bar{x}$
- ❖ Median
 - Sort the values and find the middle value (or average the two middle values)
- ❖ Why do we need two measures of centrality?
 - Skew: our variable distribution may be "lopsided"

Centrality



Position

- ❖ The median is the middle value
 - 50% of the dataset has a value less than or equal to the median

Position

- ❖ The median is the middle value
 - 50% of the dataset has a value less than or equal to the median
- ❖ A quantile generalizes the median to any %
 - Sometimes called a percentile as well
- ❖ Multiple quantiles can summarize a variable distribution
 - Common quantiles: 0%, 25%, 50%, 75%, 100%

Position


- ❖ The median is the middle value
 - 50% of the dataset has a value less than or equal to the median
- ❖ A quantile generalizes the median to any %
 - Sometimes called a percentile as well
- ❖ Multiple quantiles can summarize a variable distribution
 - Common quantiles: 0%, 25%, 50%, 75%, 100%

0% quantile is min

100% quantile is max

Spread

- ❖ Variance is a measure of how far the values in the dataset are from the mean




Low Variance

High Variance


Spread

- ❖ Variance is a measure of how far the values in the dataset are from the mean
- ❖ How to calculate it:

$$Var(x) = \frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})^2$$



Low Variance




High Variance

Spread


- ❖ Variance is a measure of how far the values in the dataset are from the mean
- ❖ How to calculate it:

$$\text{Var}(x) = \frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})^2$$

- ❖ Standard deviation is the square root of the variance
 - Same units as the variable
 - Typical distance of value from mean



Low Variance



High Variance


Limitations

- ❖ Statistics usually only capture part of the story

Limitations

- ❖ Statistics usually only capture part of the story

N = 146 ; X mean = 53.7355 ; X SD = 15.7801 ; Y mean = 49.0995 ; Y SD = 24.0605 ; Pearson correlation = -0.143



Next Time

- ❖ Read DS Ch. 3
- ❖ Finish up descriptive statistics
- ❖ Start on data wrangling